

# PostgreSQL Backup & Restore with Google Drive (rclone)

This document describes the **complete, production-safe process** to back up a PostgreSQL database from a **remote / free-tier DB instance** to **Google Drive** using **rclone**, and how to **restore it correctly** on a new server or Docker container.

It also documents **common mistakes, why they happened**, and the **correct commands**.

---

## 1. Architecture Overview

- **Source:** Remote PostgreSQL DB-only instance
- **Backup type:** Logical backup (`pg_dump -Fc`)
- **Storage:** Google Drive (via rclone)
- **Automation:** cron (daily)
- **Restore targets:**
  - New VM
  - Docker PostgreSQL container

Why logical backups? - Portable across OS, VM, Docker - Small size - Easy to restore partially or fully

---

## 2. Prerequisites

### On DB Server

- PostgreSQL installed and running
- `pg_dump` available
- Internet access

### On Local Machine (one-time)

- Google account
- Google Cloud Console access

---

## 3. Google Drive Setup (Headless / Remote Server)

### 3.1 Create Google API Credentials

1. Open Google Cloud Console
2. Create a new project (or reuse one)
3. Enable **Google Drive API**
4. Create **OAuth Client ID**
5. Type: **Desktop App**
6. Save:

- 
7.  CLIENT\_ID
  8.  CLIENT\_SECRET
- 

## 4. Install & Configure rclone

### 4.1 Install rclone

```
sudo apt update  
sudo apt install -y rclone
```

---

### 4.2 Configure rclone (IMPORTANT: Headless Mode)

```
rclone config
```

---

Selections:

```
n) New remote  
name> gdrive  
Storage> drive  
client_id> YOUR_CLIENT_ID  
client_secret> YOUR_CLIENT_SECRET  
scope> 1  
root_folder_id> (Enter)  
service_account_file> (Enter)  
Edit advanced config?> n  
Use auto config?> n    # VERY IMPORTANT
```

---

At this point **rclone may behave in one of two valid ways** depending on version:

---

#### Flow A (URL shown directly on server)

- rclone prints a long Google OAuth URL
  - You open that URL on your local browser
  - Approve access
  - Google gives a verification code
  - Paste the code back into the server terminal
- 

#### Flow B (Local-machine command — what we actually used)

- rclone prints a **command to run on your local machine**, for example:

```
rcclone authorize "drive" "CLIENT_ID" "CLIENT_SECRET"
```

Steps: 1. Run that command **on your local laptop/desktop** 2. A browser window opens automatically 3. Login to Google and approve access 4. The local CLI outputs a **JSON object containing tokens** 5. Copy that full JSON output 6. Paste it back into the **server rclone prompt** when asked

This method is **official, secure, and recommended** for fully headless servers.

---

After either flow completes, verify:

```
rcclone lsd gdrive:
```

#### 4.3 Create Backup Folder on Drive

```
rcclone mkdir gdrive:postgres-backups
```

### 5. Prepare Local Backup Directory (DB Server)

```
sudo mkdir -p /var/backups/postgres  
sudo chown postgres:postgres /var/backups/postgres
```

### 6. Backup Script (Correct & Final)

Create script:

```
sudo -u postgres nano /var/backups/postgres/daily_backup.sh
```

```
#!/bin/bash  
set -euo pipefail  
cd /  
  
DB_NAME="app_prod"  
BACKUP_DIR="/var/backups/postgres"  
DATE=$(date +%F)  
FILE="${BACKUP_DIR}/${DB_NAME}_${DATE}.dump"  
  
# Create logical backup  
pg_dump -Fc "$DB_NAME" > "$FILE"
```

```
# Upload to Google Drive
rclone copy "$FILE" gdrive:postgres-backups --quiet

# Keep only last 14 days locally
find "$BACKUP_DIR" -type f -name "*.dump" -mtime +14 -delete
```

```
sudo chmod +x /var/backups/postgres/daily_backup.sh
```

## 7. rclone Access for postgres User (CRITICAL)

```
sudo mkdir -p /var/lib/postgresql/.config/rclone
sudo cp ~/.config/rclone/rclone.conf /var/lib/postgresql/.config/rclone/
sudo chown -R postgres:postgres /var/lib/postgresql/.config
```

Test:

```
sudo -u postgres rclone lsd gdrive:
```

## 8. Manual Test (MUST DO)

```
sudo -u postgres /var/backups/postgres/daily_backup.sh
```

Verify:

```
ls -lh /var/backups/postgres
rclone ls gdrive:postgres-backups
```

## 9. Automate with cron

```
sudo -u postgres crontab -e
```

```
0 2 * * * /var/backups/postgres/daily_backup.sh >> /var/backups/postgres/
backup.log 2>&1
```

## 10. Restore Process (Docker PostgreSQL)

### 10.1 Copy Dump into Container

```
docker cp app_prod_2026-02-07.dump postgres:/tmp/app_prod.dump
```

---

### 10.2 Use the CORRECT User

Docker PostgreSQL does NOT always have `postgres` role. Use the user defined by `POSTGRES_USER`.

Example:

```
docker exec -it postgres psql -U gasops
```

---

### 10.3 Drop & Recreate Target DB (Recommended)

```
docker exec -it postgres psql -U gasops -d postgres -c "DROP DATABASE app_prod_restore;"  
docker exec -it postgres psql -U gasops -d postgres -c "CREATE DATABASE app_prod_restore;"
```

---

### 10.4 Restore Command (Correct)

```
docker exec -i postgres pg_restore  
-U gasops  
-d app_prod_restore  
--no-owner  
--no-privileges  
< app_prod_2026-02-07.dump
```

---

## 11. Verification After Restore

```
\dn  
\dt auth.*  
\dt public.*  
SELECT COUNT(*) FROM public.songs;  
SELECT COUNT(*) FROM auth.users;
```

---

## 12. Mistakes Encountered & Corrections

### ✗ Mistake 1: Using auto-config on remote server

- **Problem:** No browser available
  - **Fix:** Use manual OAuth (`Use auto config? n`)
- 

### ✗ Mistake 2: Missing backup directory

- **Problem:** nano could not save file
- **Fix:**

```
mkdir -p /var/backups/postgres
```

---

### ✗ Mistake 3: `find` permission warning

- **Cause:** postgres user cannot access `/home/ubuntu`
- **Fix:**

```
cd /
```

---

### ✗ Mistake 4: Assuming `postgres` role exists in Docker

- **Reality:** Docker creates ONLY `POSTGRES_USER`
  - **Fix:** Always restore using that user
- 

### ✗ Mistake 5: Schema already exists during restore

- **Cause:** Restoring into non-empty DB
  - **Fix:** Drop & recreate DB before restore
- 

## 13. Best Practices (Final)

- Always test restore
  - Always store backups **outside the VM**
  - Always use `--no-owner --no-privileges` across environments
  - Prefer logical backups for small/medium DBs
  - Keep backups boring and predictable
-

## 14. Outcome

This setup provides:

- Daily automated backups
- Off-server safety
- Zero cost
- Cross-environment restore
- Verified disaster recovery

This is **production-grade**, even on free tier.